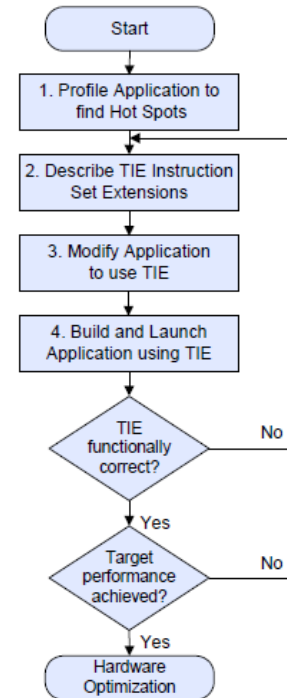


Introduction

The HW/SW-Codesign Lab is intended to give a first practical access to this topic and experiencing the potential of the conjoint design of hardware and software. For this purpose the reconfigurable processor flow from the company **Tensilica** is used (shown right). The tool suite provides convenient means for analyzing and optimizing the performance of the software. Especially the possibility to easily extend the processor's hardware architecture by customer specific instructions using the **Tensilica Instruction Extension (TIE)** language is well suited to demonstrate the interaction between hardware and software design.

Within the scope of this Lab the following 2 exercises should be done:

1. Implementation of a FIR filter design including performance analysis; HW/SW optimization: e.g. by Fusion, SIMD extension, etc.; also consider different implementation alternatives for the FIR
2. Improve the performance of a given FFT/IFFT algorithm by using basic instruction extension concepts (Fusion/SIMD/FLIX/etc.). Also consider different implementations such as the Decimation in Time (DIT) and Decimation in Frequency (DIF) algorithm.



*Tensilica® , Instruction Extension (TIE)
Language, User's Guide, 02/2014, p.4*

The written report should only include the results of the second task and need not exceed 3 pages (source code excerpts not included)!

The detailed project description can be found at the lab webpage:

<http://mns.ifn.et.tu-dresden.de/Teaching/Courses/Pages/P-HWSW-Codesign.aspx>

Author(s)

Name	E-Mail
Jörg Thalheim	joerg@higgsboson.tk
Patrick Schöps	stoepzel@stoepzel.net
Alfred Krohmer	alfred.krohmer@tu-dresden.de

If desired, the names/email will be removed from the document before publishing it on our web page.

FFT/IFFT algorithm acceleration

Source code C files (only excerpts from the parts that have been changed; please add line numbers to the code and highlight important parts):

```

dit-flix.h:
1 #ifndef FFT_ASM_H
2 #define FFT_ASM_H
3
4 #define FLIX
5
6 #ifdef FLIX
7 #define FFT_FLIX_SIMD_STORE( fr, fi, i, simd_r, simd_i) \
8     asm("{\n" \
9         "    fft_simd_store %1, %0, %2\n" \
10        "    nop\n" \
11        "    fft_simd_store %3, %0, %4\n" \
12        "}" \
13        :: "r" (i), "r" (fr), "r" (simd_r), "r" (fi), "r" (simd_i));
14
15 #define FLIX_S16I( p1, _v1, _p2, _v2) \
16     asm("{\n" \
17         "    s16i %1, %0, 0\n" \
18         "    nop\n" \
19         "    s16i %3, %2, 0\n" \
20         "}" \
21         :: "r" (p1), "r" (v1), "r" (p2), "r" (v2) \
22         : "memory");
23
24 #define FFT_FLIX_SIMD_THIRD(i, simd_r, simd_r2, simd_i, simd_i2, shift, inverse) \
25     do { \
26         /* also i += 4 in the next cycle */ \
27         asm("{\n" \
28             "    fft_simd_third %0, %1, %2, %3, %4\n" \
29             "    addi %0, %0, 4\n" \
30             "    nop\n" \
31             "}" \
32             :: "r" (i), "r" (simd_r), "r" (simd_i), "r" (shift), "r" (inverse)); \
33         /* also i -= 4 in the next cycle */ \
34         asm("{\n" \
35             "    fft_simd_third %0, %1, %2, %3, %4\n" \
36             "    addi %0, %0, -4\n" \
37             "    nop\n" \
38             "}" \
39             :: "r" (i), "r" (simd_r2), "r" (simd_i2), "r" (shift), "r" (inverse)); \
40     } while(0)
41 #define FFT_FLIX_SIMD_STORE_SHUFFLE(i, j, fr, simd_r, simd_r2, fi, simd_i, simd_i2) \
42     do { \
43         /* also i++ */ \
44         asm("{\n" \
45             "    fft_simd_store_shuffle %1, %0, %2, %5, 0\n" \
46             "    addi %0, %0, 8\n" \
47             "    fft_simd_store_shuffle %3, %0, %4, %6, 0\n" \
48             "}" \
49             /* 0      1      2      3      4 */ \
50             : "=r" (i) : "r" (fr), "r" (simd_r), "r" (fi), "r" (simd_i), \
51             "r" (simd_r2), "r" (simd_i2)); \
52         /* 5      6 */ \
53         /* also j = i + 1; */ \
54         asm("{\n" \
55             "    fft_simd_store_shuffle %3, %0, %4, %7, 1\n" \
56             "    addi %0, %0, %9\n" \
57             "    fft_simd_store_shuffle %5, %0, %6, %8, 1\n" \
58             "}" \
59             /* 0      1      2      3      4 */ \
60             : "=r" (j) : "r" (j), "r" (i), "r" (fr), "r" (simd_r), \
61             "r" (fi), "r" (simd_i), "r" (simd_r2), "r" (simd_i2), "r" (1)); \
62             /* 5      6      7      8      9 */ \
63     } while(0)
64 #else
65

```

```

66 #define FFT_FLIX_SIMD_STORE(fr, fi, i, simd_r, simd_i) \
67     FFT_SIMD_STORE(fr, i, simd_r); \
68     FFT_SIMD_STORE(fi, i, simd_i); \
69
70 #define FFT_FLIX_SIMD_THIRD(i, simd_r, simd_r2, simd_i, simd_i2, shift, inverse) \
71     FFT_SIMD_THIRD(i, simd_r, simd_i, shift, inverse); \
72     FFT_SIMD_THIRD(i + 4, simd_r2, simd_i2, shift, inverse);
73 #define FLIX_S16I(p1, v1, p2, v2) \
74     do { \
75         *(p1) = (v1); \
76         *(p2) = (v2); \
77     } while(0)
78 #define FFT_FLIX_SIMD_STORE_SHUFFLE(i, j, fr, simd_r, simd_r2, fi, simd_i, simd_i2) \
79     do { \
80         FFT_SIMD_STORE_SHUFFLE(fr, i, simd_r, simd_r2, 0); \
81         FFT_SIMD_STORE_SHUFFLE(fi, i, simd_i, simd_i2, 0); \
82         FFT_SIMD_STORE_SHUFFLE(fr, j, simd_r, simd_r2, 1); \
83         FFT_SIMD_STORE_SHUFFLE(fi, j, simd_i, simd_i2, 1); \
84         i++; \
85         j = i + 1; \
86     } while(0)
87
88 #endif // FLIX
89
90 #define FFT_FLIX_SIMD_LOAD_SHUFFLE(i, fr, simd_r, simd_r2, fi, simd_i, simd_i2, high) \
91     do { \
92         FFT_SIMD_LOAD_SHUFFLE(fr, i, simd_r, simd_r2, high); \
93         FFT_SIMD_LOAD_SHUFFLE(fi, i, simd_i, simd_i2, high); \
94     } while(0)
95
96 #endif // FFT_ASM_H

```

dit.c:

```

1 #include "dit.h"
2 #include "dit-flix.h"
3
4 #include <xtensa/tie/dit.h>
5
6
7 int fix_dit_fft(fixed fr[], fixed fi[], int size, int inverse)
8 {
9     int i, j, l, k, istep, n, m;
10    xtbool shift, inverse = _inverse;
11    register FFT_REG_SIMD simd_r, simd_i, simd_r2, simd_i2;
12
13    //number of input data
14    n = 1 << size;
15    if(n > N_WAVE) return -1;
16
17    int scale = 0;
18
19    /* decimation in time - re-order data */
20    for (m = 1; m < n; m++) {
21        int mr = FFT_BIT_REVERSE(m, size);
22
23        if(mr <= m) continue;
24        int ti = fi[m];
25        int tr = fr[m];
26
27        FLIX_S16I(&fr[m], fr[mr], &fi[m], fi[mr]);
28        FLIX_S16I(&fi[mr], ti, &fr[mr], tr);
29    }
30
31    l = 1;
32    k = LOG2_N_WAVE-1;
33    while (l < n)
34    {
35        if (inverse)
36        {
37            /* variable scaling, depending upon data */
38            shift = 0;
39            for (i = 0; i < (n / 8); i += 8)
40            {
41                if (FFT_SHIFT_CHECK(fr, i) | FFT_SHIFT_CHECK(fi, i))
42                {
43                    shift = 1;
44                    ++scale;
45                    break;
46                }

```

```

47     }
48   }
49   else
50   {
51     /* fixed scaling, for proper normalization -
52        there will be log2(n) passes, so this
53        results in an overall factor of 1/n,
54        distributed to maximize arithmetic accuracy. */
55     shift = 1;
56   }
57
58   /* it may not be obvious, but the shift will be performed
59      on each data point exactly once, during this pass. */
60   istep = 1 << 1; //step width of current butterfly
61
62   switch(l)
63   {
64     case 1:
65       for (i = 0; i < n; i += 8)
66       {
67         simd_r = FFT_SIMD_LOAD(fr, i);
68         simd_i = FFT_SIMD_LOAD(fi, i);
69         FFT_SIMD_FIRST(simd_r, simd_i, shift);
70         FFT_FLIX_SIMD_STORE(fr, fi, i, simd_r, simd_i);
71       }
72       break;
73
74     case 2:
75       for (i = 0; i < n; i += 8)
76       {
77         simd_r = FFT_SIMD_LOAD(fr, i);
78         simd_i = FFT_SIMD_LOAD(fi, i);
79         FFT_SIMD_SECOND(simd_r, simd_i, shift, inverse);
80         FFT_FLIX_SIMD_STORE(fr, fi, i, simd_r, simd_i);
81       }
82       break;
83
84     case 4:
85       WUR_FFT_SIMD_K(7);
86       for (i = 0; i < n; i += 8)
87       {
88         simd_r = FFT_SIMD_LOAD(fr, i);
89         simd_i = FFT_SIMD_LOAD(fi, i);
90         FFT_SIMD_THIRD(i, simd_r, simd_i, shift, inverse);
91         FFT_FLIX_SIMD_STORE(fr, fi, i, simd_r, simd_i);
92       }
93       break;
94
95     default:
96       WUR_FFT_SIMD_K(k);
97       for (m = 0; m < n; m += istep)
98       {
99         j = m + 1;
100        for (i = m; i < m + 1;)
101        {
102          FFT_FLIX_SIMD_LOAD_SHUFFLE(i, fr, simd_r, simd_r2, fi, simd_i, simd_i2, 0);
103          FFT_FLIX_SIMD_LOAD_SHUFFLE(j, fr, simd_r, simd_r2, fi, simd_i, simd_i2, 1);
104
105          FFT_FLIX_SIMD_THIRD(i, simd_r, simd_r2, simd_i, simd_i2, shift, inverse);
106
107          // inlines also i++ and j = i + 1
108          FFT_FLIX_SIMD_STORE_SHUFFLE(i, j, fr, simd_r, simd_r2, fi, simd_i, simd_i2);
109        }
110        break;
111      }
112   }
113   --k;
114   l = istep;
115 }
116
117 return scale;
118 }

```

Source code TIE files (please add line numbers to the code):

```

1 table SIN_WAVE 16 257 {
2   0, 201, 402, 603, 804, 1005, 1206, 1406,
3   1607, 1808, 2009, 2209, 2410, 2610, 2811, 3011,
4   3211, 3411, 3611, 3811, 4011, 4210, 4409, 4608,
5   4807, 5006, 5205, 5403, 5601, 5799, 5997, 6195,
6   6392, 6589, 6786, 6982, 7179, 7375, 7571, 7766,
7   7961, 8156, 8351, 8545, 8739, 8932, 9126, 9319,
8   9511, 9703, 9895, 10087, 10278, 10469, 10659, 10849,
9   11038, 11227, 11416, 11604, 11792, 11980, 12166, 12353,
10  12539, 12724, 12909, 13094, 13278, 13462, 13645, 13827,
11  14009, 14191, 14372, 14552, 14732, 14911, 15090, 15268,
12  15446, 15623, 15799, 15975, 16150, 16325, 16499, 16672,
13  16845, 17017, 17189, 17360, 17530, 17699, 17868, 18036,
14  18204, 18371, 18537, 18702, 18867, 19031, 19194, 19357,
15  19519, 19680, 19840, 20000, 20159, 20317, 20474, 20631,
16  20787, 20942, 21096, 21249, 21402, 21554, 21705, 21855,
17  22004, 22153, 22301, 22448, 22594, 22739, 22883, 23027,
18  23169, 23311, 23452, 23592, 23731, 23869, 24006, 24143,
19  24278, 24413, 24546, 24679, 24811, 24942, 25072, 25201,
20  25329, 25456, 25582, 25707, 25831, 25954, 26077, 26198,
21  26318, 26437, 26556, 26673, 26789, 26905, 27019, 27132,
22  27244, 27355, 27466, 27575, 27683, 27790, 27896, 28001,
23  28105, 28208, 28309, 28410, 28510, 28608, 28706, 28802,
24  28897, 28992, 29085, 29177, 29268, 29358, 29446, 29534,
25  29621, 29706, 29790, 29873, 29955, 30036, 30116, 30195,
26  30272, 30349, 30424, 30498, 30571, 30643, 30713, 30783,
27  30851, 30918, 30984, 31049, 31113, 31175, 31236, 31297,
28  31356, 31413, 31470, 31525, 31580, 31633, 31684, 31735,
29  31785, 31833, 31880, 31926, 31970, 32014, 32056, 32097,
30  32137, 32176, 32213, 32249, 32284, 32318, 32350, 32382,
31  32412, 32441, 32468, 32495, 32520, 32544, 32567, 32588,
32  32609, 32628, 32646, 32662, 32678, 32692, 32705, 32717,
33  32727, 32736, 32744, 32751, 32757, 32761, 32764, 32766,
34  32767
35 }
36
37 regfile FFT_REG_SIMD 128 4 fftsv
38
39 function [31:0] FFT VAR SHIFT([31:0] data, [3:0] sh)
40 {
41   assign FFT_VAR_SHIFT = TIEmux(sh,
42     data[31:0],
43     {data[30:0], 1'b0},
44     {data[29:0], 2'b0},
45     {data[28:0], 3'b0},
46     {data[27:0], 4'b0},
47     {data[26:0], 5'b0},
48     {data[25:0], 6'b0},
49     {data[24:0], 7'b0},
50     {data[23:0], 8'b0},
51     {data[22:0], 9'b0},
52     {data[21:0], 10'b0},
53     {data[20:0], 11'b0},
54     {data[19:0], 12'b0},
55     {data[18:0], 13'b0},
56     {data[17:0], 14'b0},
57     {data[16:0], 15'b0});
58 }
59
60 function [31:0] ADD32([31:0] a, [15:0] b) slot_shared
61 {
62   assign ADD32 = TIEadd(a, b, 1'b0);
63 }
64
65 operation FFT SIMD LOAD {in AR *base, in AR offset, out FFT REG SIMD data} {out VAddr, in MemDataIn128}
66 {
67   assign VAddr = ADD32(base, {offset[30:0], 1'b0});
68
69   wire [15:0] o1 = MemDataIn128[15:0];
70   wire [15:0] o2 = MemDataIn128[31:16];
71   wire [15:0] o3 = MemDataIn128[47:32];
72   wire [15:0] o4 = MemDataIn128[63:48];
73   wire [15:0] o5 = MemDataIn128[79:64];
74   wire [15:0] o6 = MemDataIn128[95:80];
75   wire [15:0] o7 = MemDataIn128[111:96];
76   wire [15:0] o8 = MemDataIn128[127:112];
77
78   assign data = {o1, o2, o3, o4, o5, o6, o7, o8 };

```

```

79 }
80
81 operation FFT SIMD STORE {in AR *base, in AR offset, in FFT_REG_SIMD data} {out VAddr, out MemDataOut128}
82 {
83     assign VAddr = ADD32(base, {offset[30:0], 1'b0});
84
85     wire [15:0] o1 = data[15:0];
86     wire [15:0] o2 = data[31:16];
87     wire [15:0] o3 = data[47:32];
88     wire [15:0] o4 = data[63:48];
89     wire [15:0] o5 = data[79:64];
90     wire [15:0] o6 = data[95:80];
91     wire [15:0] o7 = data[111:96];
92     wire [15:0] o8 = data[127:112];
93
94     assign MemDataOut128 = {o1, o2, o3, o4, o5, o6, o7, o8 };
95 }
96
97 immediate_range offset 0 1 1
98
99 operation FFT_SIMD_LOAD_SHUFFLE {in AR *base, in AR offset, inout FFT_REG_SIMD d1, inout FFT_REG_SIMD d2,
in offset high} {out VAddr, in MemDataIn128}
100 {
101     assign VAddr = ADD32(base, {offset[30:0], 1'b0});
102
103     wire [15:0] o1 = MemDataIn128[15:0];
104     wire [15:0] o2 = MemDataIn128[31:16];
105     wire [15:0] o3 = MemDataIn128[47:32];
106     wire [15:0] o4 = MemDataIn128[63:48];
107     wire [15:0] o5 = MemDataIn128[79:64];
108     wire [15:0] o6 = MemDataIn128[95:80];
109     wire [15:0] o7 = MemDataIn128[111:96];
110     wire [15:0] o8 = MemDataIn128[127:112];
111
112     assign d1 = TIEmux(high[0], {o1, o2, o3, o4, d1[63:0]}, {d1[127:64], o1, o2, o3, o4});
113     assign d2 = TIEmux(high[0], {o5, o6, o7, o8, d2[63:0]}, {d2[127:64], o5, o6, o7, o8});
114 }
115
116 operation FFT_SIMD_STORE_SHUFFLE {in AR *base, in AR offset, in FFT_REG_SIMD d1, in FFT_REG_SIMD d2, in
offset high} {out VAddr, out MemDataOut128}
117 {
118     assign VAddr = ADD32(base, {offset[30:0], 1'b0});
119
120     wire [15:0] o1 = TIEmux(high[0], d1[79:64], d1[15:0]);
121     wire [15:0] o2 = TIEmux(high[0], d1[95:80], d1[31:16]);
122     wire [15:0] o3 = TIEmux(high[0], d1[111:96], d1[47:32]);
123     wire [15:0] o4 = TIEmux(high[0], d1[127:112], d1[63:48]);
124     wire [15:0] o5 = TIEmux(high[0], d2[79:64], d2[15:0]);
125     wire [15:0] o6 = TIEmux(high[0], d2[95:80], d2[31:16]);
126     wire [15:0] o7 = TIEmux(high[0], d2[111:96], d2[47:32]);
127     wire [15:0] o8 = TIEmux(high[0], d2[127:112], d2[63:48]);
128
129     assign MemDataOut128 = {o5, o6, o7, o8, o1, o2, o3, o4 };
130 }
131
132 operation FFT SHIFT CHECK {in AR *addr, in AR offset, out AR needs shift} {out VAddr, in MemDataIn128}
133 {
134     assign VAddr = ADD32(addr, offset[31:1]);
135
136     wire [15:0] o1 = MemDataIn128[15:0];
137     wire [15:0] o2 = MemDataIn128[31:16];
138     wire [15:0] o3 = MemDataIn128[47:32];
139     wire [15:0] o4 = MemDataIn128[63:48];
140     wire [15:0] o5 = MemDataIn128[79:64];
141     wire [15:0] o6 = MemDataIn128[95:80];
142     wire [15:0] o7 = MemDataIn128[111:96];
143     wire [15:0] o8 = MemDataIn128[127:112];
144
145     wire s1 = (!o1[15] && o1[14]) || (o1[15] && (!o1[14] || o1[13:0] == 14'b0));
146     wire s2 = (!o2[15] && o2[14]) || (o2[15] && (!o2[14] || o2[13:0] == 14'b0));
147     wire s3 = (!o3[15] && o3[14]) || (o3[15] && (!o3[14] || o3[13:0] == 14'b0));
148     wire s4 = (!o4[15] && o4[14]) || (o4[15] && (!o4[14] || o4[13:0] == 14'b0));
149     wire s5 = (!o5[15] && o5[14]) || (o5[15] && (!o5[14] || o5[13:0] == 14'b0));
150     wire s6 = (!o6[15] && o6[14]) || (o6[15] && (!o6[14] || o6[13:0] == 14'b0));
151     wire s7 = (!o7[15] && o7[14]) || (o7[15] && (!o7[14] || o7[13:0] == 14'b0));
152     wire s8 = (!o8[15] && o8[14]) || (o8[15] && (!o8[14] || o8[13:0] == 14'b0));
153
154     assign needs_shift = {31'b0, s1 || s2 || s3 || s4 || s5 || s6 || s7 || s8 };
155 }
156
157 operation FFT BIT REVERSE {in AR m, in AR mm, out AR mr} {}

```

```

158 {
159     assign mr = {
160         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0]),
161         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1]),
162         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2]),
163         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3]),
164         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4]),
165         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4], m[5]),
166         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6]),
167         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7]),
168         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7], m[8]),
169         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7], m[8], m[9]),
170         TIEmux(mn[3:0], 1'b0, 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7], m[8], m[9], m[10]),
171         TIEmux(mn[3:0], 1'b0, 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7], m[8], m[9], m[10], m[11]),
172         TIEmux(mn[3:0], 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7], m[8], m[9], m[10], m[11], m[12]),
173         TIEmux(mn[3:0], 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7], m[8], m[9], m[10], m[11], m[12], m[13]),
174         TIEmux(mn[3:0], 1'b0, m[0], m[1], m[2], m[3], m[4], m[5], m[6], m[7], m[8], m[9], m[10], m[11], m[12], m[13], m[14])
175     };
176 }
177
178 function [15:0] SIN ([9:0] idx, [0:0] not_negate)
179 {
180     //wire [7:0] neg_idx = TIEadd(0, ~idx[7:0], 1'b1);
181     wire [8:0] x2 = TIEadd(8'h0, ~idx[7:0], 1'b1); // 512 - x
182     wire [8:0] sin_idx = TIEmux(idx[8:0] == 256 || !idx[8:8], // if x mod 512 <= 256
183         // then
184         x2[8:0],
185         // else
186         idx[8:0]
187     );
188     wire [15:0] value = SIN_WAVE[sin_idx];
189
190     assign SIN = TIEmux((idx[9:0] == 512) || !idx[9:9] ) && not negate,
191         TIEadd(16'h0, ~value[15:0], 1'b1),
192         value[15:0]);
193 }
194
195 function [31:0] FFT_TWIDDLE ([31:0] i, [1:0] x, [4:0] k, [0:0] shift, [0:0] inverse)
196 {
197     wire [31:0] inc = TIEadd(i, x, 1'b0);
198     wire [31:0] j = FFT_VAR_SHIFT(inc, k);
199
200     // 256 = N_WAVE / 4
201     wire [9:0] cos_idx = TIEadd(j, 256, 1'b0);
202     // 1023 -> 0
203     wire [15:0] sin = SIN(j, inverse);
204     wire [15:0] wr1 = SIN(cos_idx, 1);
205
206     assign FFT_TWIDDLE = {
207         TIEmux(shift, wr1, {wr1[15], wr1[15:1]}),
208         TIEmux(shift, sin, {sin[15], sin[15:1]})
209     };
210 }
211
212 function [31:0] FFT_TWIDDLE1 ([31:0] i, [1:0] x, [4:0] k, [0:0] shift, [0:0] inverse) slot_shared {
213     assign FFT_TWIDDLE1 = FFT_TWIDDLE(i, x, k, shift, inverse);
214 }
215 function [31:0] FFT_TWIDDLE2 ([31:0] i, [1:0] x, [4:0] k, [0:0] shift, [0:0] inverse) slot_shared {
216     assign FFT_TWIDDLE2 = FFT_TWIDDLE(i, x, k, shift, inverse);
217 }
218
219 function [63:0] FFT_BUTTERFLY ([63:0] data, [15:0] wr, [15:0] wi, [0:0] shift) {
220
221     // operands real parts
222     wire [15:0] r1 = data[63:48];
223     wire [15:0] r2 = data[47:32];
224
225     // operands imaginary parts
226     wire [15:0] i1 = data[31:16];
227     wire [15:0] i2 = data[15:0];
228
229     // odd real part
230     wire [31:0] oddr1 = TIEmul(wr, r2, 1'b1);
231     wire [31:0] oddr2 = TIEmul(wi, i2, 1'b1);
232     wire [15:0] oddr = TIEaddn(oddr1[30:15], ~oddr2[30:15], 16'b1);
233
234     // odd imaginary part
235     wire [31:0] oddi1 = TIEmul(wr, i2, 1'b1);
236     wire [31:0] oddi2 = TIEmul(wi, r2, 1'b1);
237     wire [15:0] oddi = TIEadd(oddi1[30:15], oddi2[30:15], 1'b0);
238
239     // even parts
240     wire [15:0] evenr = TIEmux(shift[0], r1, {r1[15], r1[15:1]});

```



```

241 wire [15:0] eveni = TIEmux(shift[0], i1, {i1[15], i1[15:1]});
242
243 // final result
244 wire [15:0] resr1 = TIEadd(evenr, oddr, 1'b0);
245 wire [15:0] resr2 = TIEadd(evenr, ~oddr, 1'b1);
246 wire [15:0] resi1 = TIEadd(eveni, oddi, 1'b0);
247 wire [15:0] resi2 = TIEadd(eveni, ~oddi, 1'b1);
248
249 assign FFT_BUTTERFLY = { resr1, resr2, resi1, resi2 };
250 }
251 function [63:0] FFT_BUTTERFLY1 ([63:0] data, [15:0] wr, [15:0] wi, [0:0] shift) slot_shared {
252     assign FFT_BUTTERFLY1 = FFT_BUTTERFLY(data, wr, wi, shift);
253 }
254 function [63:0] FFT_BUTTERFLY2 ([63:0] data, [15:0] wr, [15:0] wi, [0:0] shift) slot_shared {
255     assign FFT_BUTTERFLY2 = FFT_BUTTERFLY(data, wr, wi, shift);
256 }
257
258 // first stage: four non-interleaved butterflies
259 operation FFT_SIMD_FIRST {inout FFT_REG_SIMD fr, inout FFT_REG_SIMD fi, in BR shift} {}
260 {
261     // 1
262     wire [15:0] wr = TIEmux(shift, 16'h7fff, 16'h3fff);
263     wire [15:0] wi = 16'b0;
264
265     wire [63:0] res1 = FFT_BUTTERFLY1({fr[127:96], fi[127:96]}, wr, wi, shift);
266     wire [63:0] res2 = FFT_BUTTERFLY1({fr[95:64], fi[95:64]}, wr, wi, shift);
267     wire [63:0] res3 = FFT_BUTTERFLY2({fr[63:32], fi[63:32]}, wr, wi, shift);
268     wire [63:0] res4 = FFT_BUTTERFLY2({fr[31:0], fi[31:0]}, wr, wi, shift);
269
270     assign fr = { res1[63:32], res2[63:32], res3[63:32], res4[63:32] };
271     assign fi = { res1[31:0], res2[31:0], res3[31:0], res4[31:0] };
272 }
273
274 schedule FFT_SIMD_FIRST_schedule {FFT_SIMD_FIRST}
275 {
276     def fr 2;
277     def fi 2;
278 }
279
280 // second stage: four butterflies, two interleaved at a time
281 operation FFT_SIMD_SECOND {inout FFT_REG_SIMD fr, inout FFT_REG_SIMD fi, in BR shift, in BR inverse} {}
282 {
283     // 1
284     wire [15:0] wr1 = TIEmux(shift, 16'h7fff, 16'h3fff);
285     wire [15:0] wi1 = 16'b0;
286
287     // -j
288     wire [15:0] wr2 = 16'b0;
289     wire [15:0] wi2 = TIEmux({inverse, shift}, 16'h8001, 16'hc000, 16'h7fff, 16'h3fff);
290
291     wire [63:0] res1 = FFT_BUTTERFLY1({fr[127:112], fr[95:80], fi[127:112], fi[95:80]}, wr1, wi1, shift);
292     wire [63:0] res2 = FFT_BUTTERFLY1({fr[111:96], fr[79:64], fi[111:96], fi[79:64]}, wr2, wi2, shift);
293     wire [63:0] res3 = FFT_BUTTERFLY2({fr[63:48], fr[31:16], fi[63:48], fi[31:16]}, wr1, wi1, shift);
294     wire [63:0] res4 = FFT_BUTTERFLY2({fr[47:32], fr[15:0], fi[47:32], fi[15:0]}, wr2, wi2, shift);
295
296     assign fr = { res1[63:48], res2[63:48], res1[47:32], res2[47:32], res3[63:48], res4[63:48], res3[47:32], res4[47:32] };
297     assign fi = { res1[31:16], res2[31:16], res1[15:0], res2[15:0], res3[31:16], res4[31:16], res3[15:0], res4[15:0] };
298 }
299
300 schedule FFT_SIMD_SECOND_schedule {FFT_SIMD_SECOND}
301 {
302     def fr 2;
303     def fi 2;
304 }
305
306 function [31:0] FFT_INC_SHIFT ([31:0] i, [1:0] x, [4:0] k)
307 {
308     wire [31:0] inc = TIEadd(i, x, 1'b0);
309     assign FFT_INC_SHIFT = FFT_VAR_SHIFT(inc, k);
310 }
311
312 state FFT_SIMD_K 5 add_read_write
313
314 // third stage: four interleaved butterflies
315 operation FFT_SIMD_THIRD {in AR i, inout FFT_REG_SIMD fr, inout FFT_REG_SIMD fi, in BR shift, in BR inverse} {in FFT_SIMD_K}
316 {
317     wire [4:0] k = FFT_SIMD_K;
318
319     wire [31:0] tw1 = FFT_TWIDDLE1(i, 0, k, shift, inverse);
320     wire [15:0] wr1 = tw1[31:16];
321     wire [15:0] wi1 = tw1[15:0];
322

```

```

323 wire [31:0] tw2 = FFT_TWIDDLE1(i, 1, k, shift, inverse);
324 wire [15:0] wr2 = tw2[31:16];
325 wire [15:0] wi2 = tw2[15:0];
326
327 wire [31:0] tw3 = FFT_TWIDDLE2(i, 2, k, shift, inverse);
328 wire [15:0] wr3 = tw3[31:16];
329 wire [15:0] wi3 = tw3[15:0];
330
331 wire [31:0] tw4 = FFT_TWIDDLE2(i, 3, k, shift, inverse);
332 wire [15:0] wr4 = tw4[31:16];
333 wire [15:0] wi4 = tw4[15:0];
334
335 wire [63:0] res1 = FFT_BUTTERFLY1({fr[127:112], fr[63:48], fi[127:112], fi[63:48]}, wr1, wi1, shift);
336 wire [63:0] res2 = FFT_BUTTERFLY1({fr[111:96], fr[47:32], fi[111:96], fi[47:32]}, wr2, wi2, shift);
337 wire [63:0] res3 = FFT_BUTTERFLY2({fr[95:80], fr[31:16], fi[95:80], fi[31:16]}, wr3, wi3, shift);
338 wire [63:0] res4 = FFT_BUTTERFLY2({fr[79:64], fr[15:0], fi[79:64], fi[15:0]}, wr4, wi4, shift);
339
340 assign fr = { res1[63:48], res2[63:48], res3[63:48], res4[63:48], res1[47:32], res2[47:32], res3[47:32], res4[47:32] };
341 assign fi = { res1[31:16], res2[31:16], res3[31:16], res4[31:16], res1[15:0], res2[15:0], res3[15:0], res4[15:0] };
342 }
343
344 schedule FFT_SIMD_THIRD_schedule (FFT_SIMD_THIRD)
345 {
346     def fr 2;
347     def fi 2;
348 }
349
350
351 //-----
352
353 ctype FFT_REG_SIMD 128 128 FFT_REG_SIMD default
354
355 immediate_range st.FFT_REG_SIMD_immed2 -256 240 16
356 immediate_range ld.FFT_REG_SIMD_immed2 -256 240 16
357 format flx64_0 64 { flx64_0_slot0, flx64_0_slot1, flx64_0_slot2 }
358
359 slot opcodes flx64_0_slot0 {
360     FFT_SHIFT_CHECK, OR, FFT_SIMD_LOAD, FFT_SIMD_STORE, FFT_SIMD_LOAD_SHUFFLE, FFT_SIMD_STORE_SHUFFLE,
    FFT_SIMD_THIRD, DIF_FFT_SIMD_THIRD,
361     // XTENSA ISA OPS - usable without area overhead
362     MOVI, SSL, NOP, SLL, BGE, J, ADDI, S32I, ADDX2, L16SI, S16I, BLT, BGEI, ORBC, ANDBC, BBCI, BEQZ.N, ORB,
    ADDI.N, MOVGEZ, SRAI, OR, BEQZ, L32I.N, BEQI, BNEI, MOVI.N, BLTI, SLLI, ADD.N, MOV.N, ADD, L32I, BNE
363 }
364 slot_opcodes flx64_0_slot1 { SSL, MOVI, ADDI, J, ADDX2, NOP, MOV.N, ADD.N }
365 slot_opcodes flx64_0_slot2 { NOP, S32I, ADDI, ADDX2, L16SI, S16I, MOVI, FFT_SHIFT_CHECK, FFT_SIMD_LOAD,
    FFT_SIMD_STORE, FFT_SIMD_LOAD_SHUFFLE, FFT_SIMD_STORE_SHUFFLE }

```

Questions:

Which acceleration technique(s) has/have been used and why?

- FLIX instructions:
 - better utilize hardware by executing multiple independent instructions in parallel
 - (parallel load / store in particular, see SIMD)
- SIMD:
 - load, calculate and store multiple values at once in a single instruction
- Pipelining:
 - reduce critical path to increase maximum clock frequency
- Hardware Lookup Table for Sine function:
 - speed up lookup

Which part of the FFT algorithm did you accelerate and why?

- Bit reversal:
 - is achieved in a single instruction, without a loop like in the C implementation
 - swapping of values:
 - store operations are not parallelized via FLIX by the compiler due to possible memory overlapping
 - parallelization enforced with inline assembler
- Comparison for value ranges:
 - parallel comparison (greater than or less than a given value) in hardware
 - SIMD by processing 8 input values in one instruction
- Load / Store:
 - to load and store 128 bits at once, extra instructions are required
 - even loading and storing of 256 bits in one instruction is achieved by using the same load / store instruction in two FLIX slots
 - address calculation (base + offset) realized in hardware
- FFT calculation:
 - twiddle factor and butterfly calculations are implemented as TIE instructions
 - SIMD by processing 8 input values (4 butterflies) at once

What is the impact on execution time (speedup)?

Cycles used for FFT and inverse FFT combined for:

- unoptimized C implementation (-O3):
 - M=3 (N=8): **3 512**
 - M=8 (N=256): **230 627**
 - M=10 (N=1024): **1 114 289**
- our optimized implementation with TIE instructions, SIMD, pipelining and FLIX (-O2, compiled with feedback optimization):
 - M=3 (N=8): **327** (Speedup: **10.7**)
 - M=8 (N=256): **7 612** (Speedup: **30.3**)
 - M=10 (N=1024): **32 050** (Speedup: **34.8**)

What is the impact on the Hardware (Area)?

Our TIE implementation required about **90 000** additional gates for either DIT or DIF. When using both DIT and DIF at the same time, only about 108 000 additional gates are required because of shared functionalities.

What is the impact on the software (new instructions, modified source code, compiler intrinsics)?

By implementing parts of the algorithm in TIE, we introduced new instructions. Also, we had to restructure the core loop of the algorithm to fit the memory layout required by the parallelized load / store instructions.

Detailed Report (min. 1 page, max. 3 pages): *Step-by-step description of each optimization step that has been tried (even if no performance increase could be achieved):*

- *What software optimization or hardware acceleration strategies you tried?*
- *Refer your explanations to the C-/TIE-code making use of the lines numbers*
- *How did the performance increase (or decrease)?*
- *What are the costs for the performance increase?*
- *Make a trade off performance vs. increased costs.*

Conclude with an overall summary: Which combination of optimization and hardware acceleration strategies do you suggest, i.e., yields the best performance results or speed/area tradeoff? Summarize the experiences you gained with this task.

Note: For the following descriptions, we used $M=8$ ($N=256$) with $-O3$ but without feedback optimization to measure the spent cycles. Measured is the FFT as well as the inverse FFT combined.

At first we ran the profiling tool from within Xtensa Xplorer and found that the innermost loop core required the most cycles. Therefore we introduced a new TIE instruction that does the actual butterfly calculation. This early version required 109 832 cycles (speedup 2.1) and 16 444 gates.

The second hot-spot we optimized was the bit reverse at the beginning of the algorithm. We implemented the bit reversal itself as a single TIE instruction. This had the advantage that it used merely one cycle instead of using a loop. Also we implemented the calculation of the twiddle factors in another TIE instruction by using a hardware lookup table for the required sine values. We achieved a total of 89 107 cycles (speedup 2.6) with using 32 793 gates.

Next we wanted to optimize loop headers throughout the C implementation by introducing corresponding TIE instructions. However the achieved speedup was merely noticeable or the program was even slower. This might be because the optimizations the C compiler tries to do are impaired by the newly introduced instructions it can't comprehend. Therefore, we abandoned this approach.

Another hot-spot was the check if any of the input values are within the specified range. We optimized this step by doing both comparisons (greater than or less than a given value) in another TIE instruction (dit.tie: 132-155). Also we introduced auto-generated FLIX optimization for the first time. This version finished after 77 361 cycles (speedup 3.0) but required 82 932 gates which is a tremendous increase for such a small gain in speedup.

However, after several manual changes of the FLIX slot assignment (dit.tie: 359-365) we managed to reduce the number of gates to 45 079 while decreasing the number of required cycles to 60 309 (speedup 3.8).

For the next step, we tried to use the TIE_{mac} function instead of TIE_{mul} and TIE_{add} in our TIE instructions. This resulted only in slightly lower gate usage and no performance gain but produced results that differed by a small factor from the C implementation, which we wanted to avoid.

Next we introduced SIMD operation for the instruction that checks for the value ranges. This reduced the needed cycles to 48 327 (speedup 4.8) and required gates to 59 697.

After that we moved the twiddle factor calculation inside the butterfly calculation which resulted in a longer computation time, requiring 53594 cycles, but also reduced the number of used gates to 51597.

Our next big improvement was the introduction of 4-fold SIMD processing of the actual butterfly calculation as well as the use of some small snippets of inline assembly code to force the compiler to do parallel store operations (see dit-flix.h).

For this optimization to work, we needed to implement separate handling of the first three stages of the FFT (dit.c: 62-112). Also we implemented custom load and store operations (dit.tie: 65-95) to be able to transfer 128 bits of memory within one instruction (and 256 bits within one cycle when using FLIX).

For stage four and beyond we implemented another load and store operation (dit.tie: 99-130) so that we could fetch enough data (in the right order) for two 4-fold butterflies within 2 cycles. By outsourcing the butterfly calculation into shared functions (dit.tie: 251-256) we could share those four required hardware instances across the separate instructions for the different stages. For stage one and two, we hardcoded the required twiddle factors. For stage three and beyond, four instances of the twiddle factor calculation were needed.

Altogether this optimization resulted in huge performance gain as the number of spent cycles dropped to 9 640 (speedup 23.9). However, this came with a trade-off regarding hardware area usage as our TIE implementation now required 175 642 gates.

As a side effect of using 4-fold SIMD our optimized algorithm requires a problem size of at least $M=3$ ($N=8$) to work properly. Because of the generally limited size of the sine table the maximum problem size is $M=10$ ($N=1024$).

Several smaller optimizations followed, mainly aimed at reducing the number of required gates.

We managed to reduce the size of the hardware lookup table for the sine values to a quarter of the original size by using the symmetry of the sine function (dit.tie: 178-193). Together with code cleanup by removing now unneeded functions we reduced the gate count to 118 511. (9284 cycles)

Finally, by using the schedule option (e.g. dit.tie: 344-348), we pipelined the butterfly calculation instructions across two clock cycles, slightly reducing the performance, but also basically cutting the critical path in half, possibly allowing for higher clock rates and reducing the required gate count. We ended up using a total of 90 503 gates.

Surprisingly, using the `-O2` instead of the `-O3` compiler option, we achieved a better performance. Together with enabled software feedback optimization, we managed to reduce the number of required clock cycles to 7612 (speedup 30.3). An even higher speedup was achieved for $M=10$ ($N=1024$); see above.

Actually, using `-O3` seems to break our current implementation as we apparently get random output values. This might be because of some optimizations the compiler tries to do on its own which interfere with our inline assembler instructions although we correctly annotated which variables we need to use and which we modify in the assembly code.

To compare DIT and DIF we also implemented DIF in C code. In contrast to DIT, we had to put the bit reverse at the end of the algorithm, reverse the loop direction and implement the modified butterfly. For optimization we reused most of the TIE instructions we implemented for DIT, we only had to rewrite the butterfly calculation for DIF.

For the difference between those two implementations we only noticed that the DIF implementation in C took more cycles to complete. When using the TIE implementation, execution time and hardware area usage were roughly the same as for DIT.

Summary:

- Although SIMD increases hardware area, it also improves performance by an even higher number, especially for bigger problem sizes.
- FLIX complicates instruction decoding and therefore also uses more area by multiplying the hardware components for instructions, but can increase the performance by parallelizing instructions in many places. Not only for custom TIE instructions. Also FLIX is the only option to use both load/store-units in one clock cycle.
- Pipelining delays single instructions resulting in lower overall performance when following instructions depend on previous results. On the other hand, it cuts the critical path and permits a higher clock rate. By sharing hardware resources between different pipeline stages, one can further reduce the hardware usage.

In this task we learned a lot about the interaction between hardware and software; how much performance certain hardware implementations yield and with which area trade-off they come. We also learned how to use profiling to detect hot-spots in given applications. For the first time we used a VLIW / FLIX platform and wrote inline assembly code.

After all we gained much experience and also had fun trying to optimize almost every bit of code to a maximum while still keeping the costs in mind.